

Bayesian Learning

Lecture 9 - Hamiltonian Monte Carlo and Variational Inference

Mattias Villani 🧑

Department of Statistics
Stockholm University



Lecture overview

- **Hamiltonian Monte Carlo**
- **Variational Inference**

Hamiltonian Monte Carlo

- When $\theta = (\theta_1, \dots, \theta_p)^\top$ is **high-dimensional**, $p(\theta|\mathbf{y})$ usually located in some subregion of \mathbb{R}^p with complicated geometry.
- MH: hard to find good proposal distribution $q(\cdot|\theta^{(i-1)})$.
- MH: use very small step sizes otherwise too many rejections.
- **Hamiltonian Monte Carlo (HMC):**
 - ▶ distant proposals **and**
 - ▶ high acceptance probabilities.
- HMC: add extra **momentum** parameters $\phi = (\phi_1, \dots, \phi_p)^\top$ and sample from

$$p(\theta, \phi|\mathbf{y}) = p(\theta|\mathbf{y}) p(\phi)$$

.

Hamiltonian Monte Carlo

- Physics: **Hamiltonian** system $H(\boldsymbol{\theta}, \boldsymbol{\phi}) = U(\boldsymbol{\theta}) + K(\boldsymbol{\phi})$, where U is the **potential energy** and K is the **kinetic energy**.

- **Hamiltonian Dynamics**

$$\begin{aligned}\frac{d\theta_i}{dt} &= \frac{\partial H}{\partial \phi_i} = \frac{\partial K}{\partial \phi_i}, \\ \frac{d\phi_i}{dt} &= -\frac{\partial H}{\partial \theta_i} = -\frac{\partial U}{\partial \theta_i}\end{aligned}$$

- Hockey puck sliding over a friction-less surface: [illustration](#).
- **Posterior sampling**: $U(\boldsymbol{\theta}) = -\log[p(\boldsymbol{\theta})p(\mathbf{y}|\boldsymbol{\theta})]$.
- Momentum: $\boldsymbol{\phi} \sim N(\mathbf{0}, \mathbf{M})$ where \mathbf{M} is the mass matrix and

$$K(\boldsymbol{\phi}) = -\log[p(\boldsymbol{\phi})] = \frac{1}{2}\boldsymbol{\phi}^\top \mathbf{M}^{-1}\boldsymbol{\phi} + \text{const}$$

- If we could propose $\boldsymbol{\theta}$ in continuous time (spoiler: we can't), the acceptance probability would be one.

Hamiltonian Monte Carlo

■ Hamiltonian Dynamics

$$\begin{aligned}\frac{d\theta_i}{dt} &= [\mathbf{M}^{-1}\boldsymbol{\phi}]_i, \\ \frac{d\phi_i}{dt} &= \frac{\partial \log p(\boldsymbol{\theta}|\mathbf{y})}{\partial \theta_i}\end{aligned}$$

approximated using L steps with the **leapfrog algorithm**

$$\begin{aligned}\phi_i\left(t + \frac{\varepsilon}{2}\right) &= \phi_i(t) + \frac{\varepsilon}{2} \frac{\partial \log p(\boldsymbol{\theta}|\mathbf{y})}{\partial \theta_i} \Big|_{\theta(t)} \\ \theta_i(t + \varepsilon) &= \theta_i(t) + \varepsilon \mathbf{M}^{-1} \phi_i\left(t + \frac{\varepsilon}{2}\right), \\ \phi_i(t + \varepsilon) &= \phi_i\left(t + \frac{\varepsilon}{2}\right) + \frac{\varepsilon}{2} \frac{\partial \log p(\boldsymbol{\theta}|\mathbf{y})}{\partial \theta_i} \Big|_{\theta(t+\varepsilon)},\end{aligned}$$

where ε is the **step size**.

■ **Discretization** \Rightarrow acceptance probability drops with ε .

The Hamiltonian Monte Carlo algorithm

■ Initialize $\theta^{(0)}$ and iterate for $i = 1, 2, \dots$

- 1 Sample the starting **momentum** $\phi_s \sim N(0, \mathbf{M})$
- 2 Simulate new values for (θ_p, ϕ_p) by iterating the **leapfrog algorithm** L times, starting in $(\theta^{(i-1)}, \phi_s)$.
- 3 Compute the **acceptance probability**

$$\alpha = \min \left(1, \frac{p(\mathbf{y}|\theta_p)p(\theta_p)}{p(\mathbf{y}|\theta^{(i-1)})p(\theta^{(i-1)})} \frac{p(\phi_p)}{p(\phi_s)} \right)$$


- 4 With probability α set $\theta^{(i)} = \theta_p$ and $\phi^{(i)} = \phi_p$ otherwise.

Tuning Hamiltonian Monte Carlo


- HMC is very efficient, but **needs careful tuning** to work.
- **Tuning parameters:**
 - ▶ **stepsize** ε ,
 - ▶ **number of leapfrog** iterations L and
 - ▶ **mass matrix** M . (hello $J_x^{-1}(\hat{\theta})$, our old friend)
- **No U-turn** sampler:
 - ▶ **Warm-up** to determine ε and L to get good acceptance rate.
 - ▶ Avoids U-turns in the Hamiltonian proposals.
- Drawbacks of HMC:
 - ▶ Need to **evaluate gradient of log posterior** many times during Hamiltonian iterations. Costly! (Subsampling HMC).
 - ▶ Difficulty with **multimodality** (true for most algorithms).
 - ▶ Standard HMC cannot handle **discrete parameters**. Mixture example. Some recent progress.

Comparing algorithms for bivariate normal


Gibbs




RWM - Identity M




HMC - Diagonal M



HMC - Full M



Comparing algorithms for bivariate normal



Variational Inference

- VI: approximate the posterior $p(\boldsymbol{\theta}|\mathbf{y})$ with a (simpler) distribution $q(\boldsymbol{\theta})$.

- Before: **Normal approximation** from optimization:

$$q(\boldsymbol{\theta}) = N\left[\tilde{\boldsymbol{\theta}}, J_{\mathbf{x}}^{-1}(\tilde{\boldsymbol{\theta}})\right]$$

- **Mean field Variational Inference (VI)**:

$$q(\boldsymbol{\theta}) = \prod_{i=1}^p q_i(\theta_i)$$

- **Parametric VI**: Parametric family $q_{\lambda}(\boldsymbol{\theta})$ with parameters λ .
- Find the $q(\boldsymbol{\theta})$ that **minimizes the Kullback-Leibler distance** between the true posterior p and the approximation q :

$$KL(q, p) = \int q(\boldsymbol{\theta}) \ln \frac{q(\boldsymbol{\theta})}{p(\boldsymbol{\theta}|\mathbf{y})} d\boldsymbol{\theta} = E_q \left[\ln \frac{q(\boldsymbol{\theta})}{p(\boldsymbol{\theta}|\mathbf{y})} \right].$$

Mean field approximation

- **Mean field VI** is based on factorized approximation:

$$q(\boldsymbol{\theta}) = \prod_{j=1}^p q_j(\theta_j)$$

- **No specific functional forms** are assumed for the $q_j(\theta)$.
- **Optimal densities** can be shown to satisfy:

$$q_j(\theta_j) \propto \exp \left(E_{-\theta_j} \ln p(\mathbf{y}, \boldsymbol{\theta}) \right)$$

where $E_{-\theta_j}(\cdot)$ is the expectation with respect to $\prod_{k \neq j} q_k(\theta_k)$.

- **Structured mean field approximation**. Group subset of parameters in tractable blocks. Similar to Gibbs sampling.

Mean field approximation - algorithm

- Initialize: $q_2^*(\theta_2), \dots, q_M^*(\theta_p)$

- Repeat until convergence:

- ▶ $q_1^*(\theta_1) \leftarrow \frac{\exp[E_{-\theta_1} \ln p(\mathbf{y}, \theta)]}{\int \exp[E_{-\theta_1} \ln p(\mathbf{y}, \theta)] d\theta_1}$

- ▶ \vdots

- ▶ $q_p^*(\theta_p) \leftarrow \frac{\exp[E_{-\theta_p} \ln p(\mathbf{y}, \theta)]}{\int \exp[E_{-\theta_p} \ln p(\mathbf{y}, \theta)] d\theta_p}$

- Note: no assumptions about parametric form of the $q_i(\theta)$.

- Optimal $q_i(\theta)$ often **turn out** to be parametric (normal etc).

- Just update hyperparameters in the optimal densities.

Mean field approximation - Normal model

- **Model:** $X_i | \theta, \sigma^2 \stackrel{\text{iid}}{\sim} N(\theta, \sigma^2)$.
- **Prior:** $\theta \sim N(\mu_0, \tau_0^2)$ **independent** of $\sigma^2 \sim \text{Inv-}\chi^2(\nu_0, \sigma_0^2)$.
- **Mean-field approximation:** $q(\theta, \sigma^2) = q_\theta(\theta) \cdot q_{\sigma^2}(\sigma^2)$.
- Optimal densities

$$q_\theta^*(\theta) \propto \exp \left[E_{q(\sigma^2)} \ln p(\theta, \sigma^2, \mathbf{x}) \right]$$
$$q_{\sigma^2}^*(\sigma^2) \propto \exp \left[E_{q(\theta)} \ln p(\theta, \sigma^2, \mathbf{x}) \right]$$

Normal model - VB algorithm

■ Variational density for σ^2

$$\sigma^2 \sim \text{Inv} - \chi^2(\tilde{\nu}_n, \tilde{\sigma}_n^2)$$

$$\text{where } \tilde{\nu}_n = \nu_0 + n \text{ and } \tilde{\sigma}_n^2 = \frac{\nu_0 \sigma_0^2 + \sum_{i=1}^n (x_i - \tilde{\mu}_n)^2 + n \cdot \tilde{\tau}_n^2}{\nu_0 + n}$$

■ Variational density for θ

$$\theta \sim N(\tilde{\mu}_n, \tilde{\tau}_n^2)$$

where

$$\tilde{\tau}_n^2 = \frac{1}{\frac{n}{\tilde{\sigma}_n^2} + \frac{1}{\tau_0^2}}$$


$$\tilde{\mu}_n = \tilde{w} \bar{x} + (1 - \tilde{w}) \mu_0,$$

where


$$\tilde{w} = \frac{\frac{n}{\tilde{\sigma}_n^2}}{\frac{n}{\tilde{\sigma}_n^2} + \frac{1}{\tau_0^2}}$$

Normal example from Murphy ($\lambda = 1/\sigma^2$)


Initial values




After updating q_μ



After updating q_{σ^2}



At convergence



Probit regression

- **Model:**

$$\Pr(y_i = 1 | \mathbf{x}_i) = \Phi(\mathbf{x}_i^T \boldsymbol{\beta})$$

- **Prior:** $\boldsymbol{\beta} \sim N(0, \Sigma_{\boldsymbol{\beta}})$. For example: $\Sigma_{\boldsymbol{\beta}} = \tau^2 I$.

- **Latent variable formulation** with $\mathbf{u} = (u_1, \dots, u_n)'$

$$\mathbf{u} | \boldsymbol{\beta} \sim N(\mathbf{X}\boldsymbol{\beta}, 1)$$

and

$$y_i = \begin{cases} 0 & \text{if } u_i \leq 0 \\ 1 & \text{if } u_i > 0 \end{cases}$$

- Factorized **variational approximation**

$$q(\mathbf{u}, \boldsymbol{\beta}) = q_{\mathbf{u}}(\mathbf{u}) q_{\boldsymbol{\beta}}(\boldsymbol{\beta})$$

VI for probit regression

■ VI posterior

$$\beta \sim N\left(\tilde{\mu}_\beta, \left(\mathbf{X}^\top \mathbf{X} + \Sigma_\beta^{-1}\right)^{-1}\right)$$

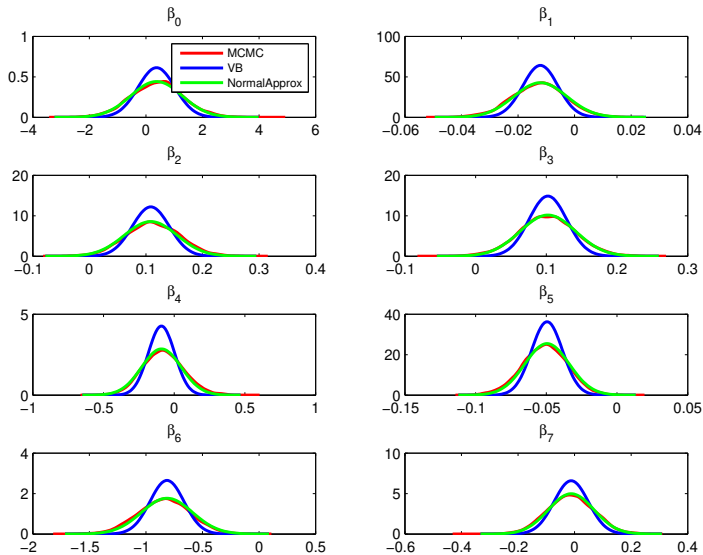
where

$$\tilde{\mu}_\beta = \left(\mathbf{X}^\top \mathbf{X} + \Sigma_\beta^{-1}\right)^{-1} \mathbf{X}^\top \tilde{\mu}_\mathbf{u}$$

and

$$\tilde{\mu}_\mathbf{u} = \mathbf{X} \tilde{\mu}_\beta + \frac{\phi(\mathbf{X} \tilde{\mu}_\beta)}{\Phi(\mathbf{X} \tilde{\mu}_\beta)^y [\Phi(\mathbf{X} \tilde{\mu}_\beta) - \mathbf{1}_n]^{1_n - y}}.$$

Probit example (n=200 observations)



Probit example

